


Hazardous knowledge. High compliance. Low cost.

Benchmark performance and access cost across two commercial API models

● Qwen 3.6* ● GLM 5.2*



* Because this is a closed API deployment, I cannot independently verify the served weights, quantization, system prompt, or all surrounding controls.

Source: author-supplied benchmark results. “Useful” refers to rubric scoring, not independently verified real-world effectiveness.